

Understanding Generative AI - Quiz Questions

Name: Kaushik Jayaprakash

Subject: INFO 7390 - Advanced Data Science

NEU ID: 002300709

Topic: Transformer-Based Text Generation Using GPT Models

The following 15 multiple-choice questions are designed to reinforce core learning objectives from my Book Chapter on **Transformer-based generative models for text generation**, focusing on GPT-style architectures, training methods, decoding strategies, and limitations. The questions cover recall, application, and deeper analysis to ensure a comprehensive understanding of generative language modeling.

SECTION 1 — Recall Questions (5 Questions)

Question 1 — Definition of a Transformer

Which component is responsible for enabling a Transformer to handle long-range dependencies in text?

- A) Convolution layers
- B) Recurrent loops
- C) Self-attention mechanism
- D) Pooling layers
- E) Batch normalization

Correct Answer: C

Explanation:

Self-attention allows the model to compare each token with every other token, capturing long-range dependencies that RNNs struggle with. Options A and D are used in CNNs; B is used in RNNs; E is not responsible for dependency modeling.

Reference: Section 2.1 “Transformer Architecture Basics,” p. 5.

Question 2 — Tokenization

GPT-style models primarily rely on which tokenization method?

- A) Word-level tokenization
- B) Byte-Pair Encoding (BPE)
- C) Character-level encoding

- D) One-hot tokenization
- E) Lemmatization-based encoding

Correct Answer: B

Explanation:

GPT models use BPE to efficiently represent rare words while keeping vocabulary size manageable. Other methods either inflate vocabulary (A) or lose semantic structure (C).

Reference: Section 2.2 “Tokenization and Embeddings,” p. 7.

Question 3 — Training Objective

What is the core training objective of a GPT-style language model?

- A) Predict masked tokens
- B) Predict the next token in a sequence
- C) Classify sentiment
- D) Minimize reconstruction loss
- E) Select a response from predefined choices

Correct Answer: B

Explanation:

GPT uses causal language modeling (next-token prediction). Option A describes BERT; D describes autoencoders; C and E are downstream tasks.

Reference: Section 3.1 “Causal Language Modeling,” p. 10.

Question 4 — Position Encoding

Why do Transformer models need positional encodings?

- A) To reduce training time
- B) To add sequential order information
- C) To decrease overfitting
- D) To improve loss scaling
- E) To reduce model parameters

Correct Answer: B

Explanation:

Transformers process all tokens in parallel, so positional encodings provide ordering information. Other options don't reflect their function.

Reference: Section 2.3 "Positional Encodings," p. 6.

Question 5 — Decoding Strategy

Which decoding method selects the top-k most probable next tokens before sampling?

- A) Greedy
- B) Beam search
- C) Top-k sampling
- D) Top-p sampling
- E) Random sampling

Correct Answer: C

Explanation:

Top-k sampling restricts sampling to the top-k candidates, controlling randomness. Top-p samples from a dynamic set; greedy picks only the max-probability token.

Reference: Section 4.1 "Decoding Methods," p. 12.

SECTION 2 — Application Questions (5 Questions)

Question 6 — Choosing a Decoding Strategy

You want the model to generate creative marketing slogans while avoiding incoherent text. Which decoding setting is most appropriate?

- A) Greedy decoding
- B) Temperature = 0.1
- C) Temperature = 1.2 + Top-k sampling
- D) Random sampling with no constraints
- E) Beam search

Correct Answer: C

Explanation:

Temperature > 1.0 increases creativity, while top-k restricts chaos. Greedy or beam search produces repetitive text; random sampling is too uncontrolled.

Reference: Section 4.2 “Balancing Creativity and Coherence,” p. 13.

Question 7 — Handling Long Prompts

A user provides a very long prompt and the model truncates part of it. Which configuration should you adjust?

- A) Learning rate
- B) Max position embeddings
- C) Batch size
- D) Tokenizer vocabulary size
- E) Dropout

Correct Answer: B

Explanation:

The max position embedding length determines how many tokens the model can attend to. Other parameters do not influence context window size.

Reference: Section 2.4 “Context Window Limitations,” p. 8.

Question 8 — Detecting Overfitting

During fine-tuning on a small dataset, the training loss decreases, but text generations become repetitive. What does this indicate?

- A) Underfitting
- B) Overfitting
- C) Poor tokenization
- D) Incorrect optimizer
- E) Low temperature

Correct Answer: B

Explanation:

Repetition in generations often reflects overfitting to narrow patterns. Low temperature could worsen repetition but doesn't explain decreasing training loss alone.

Reference: Section 5.1 "Limitations and Failure Modes," p. 16.

Question 9 — Improving Model Diversity

A model keeps generating identical responses. Which modification helps increase output diversity?

- A) Lowering top-p
- B) Lowering temperature
- C) Increasing number of training epochs
- D) Increasing temperature
- E) Switching to greedy decoding

Correct Answer: D

Explanation:

Higher temperature increases randomness, improving variety. Lowering top-p or using greedy decoding reduces diversity.

Reference: Section 4.1 "Temperature Scaling," p. 12.

Question 10 — Domain Adaptation

You fine-tune GPT-2 on legal documents, and it begins generating structured legal language. Which concept explains this shift?

- A) Zero-shot learning
- B) Catastrophic forgetting
- C) Domain adaptation
- D) Data augmentation
- E) Reinforcement learning

Correct Answer: C

Explanation:

Domain adaptation occurs when the model adjusts to patterns in a new domain. Catastrophic forgetting would imply loss of earlier ability, not a stylistic shift.

Reference: Section 3.3 “Fine-Tuning and Domain Adaptation,” p. 11.

SECTION 3 — Analysis & Evaluation Questions (5 Questions)**Question 11 — Comparing GPT and RNN Performance**

Why do GPT-style Transformers outperform LSTM-based models on long-form text generation?

- A) LSTMs store more information in their hidden states
- B) Transformers use recurrence and attention jointly
- C) Transformers capture long-range dependencies efficiently through self-attention
- D) LSTMs train faster than Transformers
- E) Transformers require no training data

Correct Answer: C

Explanation:

Self-attention lets the model access any token at any position, enabling superior handling of long context. LSTMs struggle with vanishing gradients over long distances.

Reference: Section 2.1 “Transformer Advantages Over RNNs,” p. 5.

Question 12 — Ethical Risk Assessment

A model trained on biased social media text generates toxic outputs. What is the root cause?

- A) Incorrect decoding strategy
- B) Too many layers
- C) Bias in training data
- D) Low-quality tokenizer
- E) Insufficient GPU memory

Correct Answer: C

Explanation:

Generative models reproduce biases present in their datasets. Decoding strategies or GPU memory do not introduce social biases.

Reference: Section 5.2 “Bias, Ethics, and Safety,” p. 17.

Question 13 — Perplexity Interpretation

Model A has lower perplexity than Model B on the same dataset. Which interpretation is correct?

- A) Model A memorized the dataset
- B) Model A assigns higher probability to true next tokens
- C) Model B is overfitting
- D) Model B must use a different tokenizer
- E) Model A has fewer parameters

Correct Answer: B

Explanation:

Lower perplexity means the model more accurately predicts next tokens. It does not necessarily imply overfitting or model size differences.

Reference: Section 3.4 “Evaluation Metrics,” p. 12.

Question 14 — Impact of Training Data Structure

Why does GPT-2 fine-tuned on Wikitext-2 generate more formal text than GPT-2 fine-tuned on Reddit posts?

- A) GPT-2 cannot generate informal language
- B) Model size determines writing style
- C) Transformers ignore training data style
- D) Generative models internalize statistical patterns from their datasets
- E) Reddit uses a different tokenizer

Correct Answer: D

Explanation:

Generative models learn distributional patterns, including tone and structure. Fine-tuning shifts the model toward the domain of the new dataset.

Reference: Section 3.3 “Dataset Influence on Output Style,” p. 11.

Question 15 — Evaluating Decoding Trade-offs

A company wants controllable, non-chaotic text generation while still keeping slight creativity. Which approach offers the best trade-off?

- A) Pure greedy decoding
- B) Beam search with no sampling
- C) Top-k or top-p sampling with moderate temperature
- D) Fully random sampling
- E) Temperature set to 0

Correct Answer: C

Explanation:

Top-k or top-p restrict randomness to meaningful tokens, while moderate temperature adds controlled creativity. Greedy and beam search are too rigid; random sampling is too chaotic.

Reference: Section 4.2 “Controlling Generation Quality,” p. 14.